

1. Viewer flags a subtitle segment
(translation / pinyin / text / other)

2. Threshold check:
 ≥ 3 reports AND $\geq 25\%$ of viewers

3. Context window assembled
(± 3 surrounding segments)

4. Prompt sent to Llama 3.3 70B
via Groq (JSON-mode output)

5. Timestamps re-applied server-side;
correction validated

6. Stored as a patch in PostgreSQL
(SegmentCorrection) -- MongoDB original untouched

7. Merged over original segment
at read time → shown to all viewers